

## 04\_Merged\_Data\_Pretraining

April 15, 2022

### 0.1 Predicting Airline Delays

Notebook: Data Preparation Notebook

Team: Jimmy Nguyen, Maha Jayapal, Roberto Cancel

```
[1]: !pip install --upgrade numpy #ensure numpy and pandas are upgraded to same
      ↪versions for easier exploration (avoiding errors)
      !pip install --upgrade pandas #ensure numpy and pandas are upgraded to same
      ↪versions for easier exploration (avoiding errors)

import boto3 # AWS SDK for Python
import io # for encoding issues with raw data sets
from io import StringIO # converting dataframe to csv and uploading to s3
      ↪bucket /transformed folder

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

```
/opt/conda/lib/python3.7/site-packages/secretstorage/dhcrypto.py:16:
CryptographyDeprecationWarning: int_from_bytes is deprecated, use int.from_bytes
instead
  from cryptography.utils import int_from_bytes
/opt/conda/lib/python3.7/site-packages/secretstorage/util.py:25:
CryptographyDeprecationWarning: int_from_bytes is deprecated, use int.from_bytes
instead
  from cryptography.utils import int_from_bytes
Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages
(1.21.5)
WARNING: Running pip as the 'root' user can result in broken permissions
and conflicting behaviour with the system package manager. It is recommended to
use a virtual environment instead: https://pip.pypa.io/warnings/venv
/opt/conda/lib/python3.7/site-packages/secretstorage/dhcrypto.py:16:
CryptographyDeprecationWarning: int_from_bytes is deprecated, use int.from_bytes
```

```

instead
  from cryptography.utils import int_from_bytes
/opt/conda/lib/python3.7/site-packages/secretstorage/util.py:25:
CryptographyDeprecationWarning: int_from_bytes is deprecated, use int.from_bytes
instead
  from cryptography.utils import int_from_bytes
Requirement already satisfied: pandas in /opt/conda/lib/python3.7/site-packages
(1.3.5)
Requirement already satisfied: pytz>=2017.3 in /opt/conda/lib/python3.7/site-
packages (from pandas) (2021.3)
Requirement already satisfied: python-dateutil>=2.7.3 in
/opt/conda/lib/python3.7/site-packages (from pandas) (2.8.1)
Requirement already satisfied: numpy>=1.17.3 in /opt/conda/lib/python3.7/site-
packages (from pandas) (1.21.5)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.7/site-
packages (from python-dateutil>=2.7.3->pandas) (1.14.0)
WARNING: Running pip as the 'root' user can result in broken permissions
and conflicting behaviour with the system package manager. It is recommended to
use a virtual environment instead: https://pip.pypa.io/warnings/venv

```

[2]: *# INGEST Merged DATA*

```

s3_client = boto3.client("s3")

BUCKET='ads-508-airline'
KEY='merged/Dec_modeling.csv'

response = s3_client.get_object(Bucket=BUCKET, Key=KEY)
dec_merged = pd.read_csv(response.get("Body"))
dec_merged.head()

```

```

[2]:   DAY_OF_MONTH  DAY_OF_WEEK  OP_UNIQUE_CARRIER  TAIL_NUM  ORIGIN_AIRPORT_ID  \
0           8           7           WN      N8651A           15016
1           8           7           WN      N939WN           15016
2           8           7           WN      N7741C           15016
3           8           7           WN      N550WN           15016
4           8           7           WN      N8319F           15016

   ORIGIN DEST  DEP_DEL15  DEP_TIME_BLK  ARR_TIME_BLK  ...  \
0   STL   SAN         0.0    1100-1159    1300-1359  ...
1   STL   SAT         0.0    1200-1259    1400-1459  ...
2   STL   SAT         0.0    2100-2159    0001-0559  ...
3   STL   SEA         0.0    0900-0959    1200-1259  ...
4   STL   SFO         1.0    1800-1859    2000-2059  ...

```

	CARRIER_NAME	PILOTS_COPILOTS	PASSENGER_HANDLING	\
0	Southwest Airlines Co.	8989	9668	
1	Southwest Airlines Co.	8989	9668	
2	Southwest Airlines Co.	8989	9668	
3	Southwest Airlines Co.	8989	9668	
4	Southwest Airlines Co.	8989	9668	

	PASS_GEN_SVC_ADMIN	MAINTENANCE	PRCP	SNOW	SNWD	TMAX	AWND
0	15475	2482	0.02	0.0	0.0	58.0	9.84
1	15475	2482	0.02	0.0	0.0	58.0	9.84
2	15475	2482	0.02	0.0	0.0	58.0	9.84
3	15475	2482	0.02	0.0	0.0	58.0	9.84
4	15475	2482	0.02	0.0	0.0	58.0	9.84

[5 rows x 25 columns]

```
[3]: dec_merged.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 558026 entries, 0 to 558025
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   DAY_OF_MONTH           558026 non-null int64
1   DAY_OF_WEEK            558026 non-null int64
2   OP_UNIQUE_CARRIER     558026 non-null object
3   TAIL_NUM               558026 non-null object
4   ORIGIN_AIRPORT_ID      558026 non-null int64
5   ORIGIN                 558026 non-null object
6   DEST                   558026 non-null object
7   DEP_DEL15              558026 non-null float64
8   DEP_TIME_BLK           558026 non-null object
9   ARR_TIME_BLK           558026 non-null object
10  CANCELLED               558026 non-null float64
11  CRS_ELAPSED_TIME        558026 non-null float64
12  DISTANCE                558026 non-null float64
13  DISTANCE_GROUP          558026 non-null int64
14  AIRLINE_ID              558026 non-null int64
15  CARRIER_NAME           558026 non-null object
16  PILOTS_COPILOTS         558026 non-null int64
17  PASSENGER_HANDLING      558026 non-null int64
18  PASS_GEN_SVC_ADMIN      558026 non-null int64
19  MAINTENANCE             558026 non-null int64
20  PRCP                    558026 non-null float64
21  SNOW                    558026 non-null float64
22  SNWD                    558026 non-null float64
23  TMAX                    558026 non-null float64
```

```

24  AWND                      558026 non-null  float64
dtypes: float64(9), int64(9), object(7)
memory usage: 106.4+ MB

```

```
[4]: dec_merged.shape
```

```
[4]: (558026, 25)
```

## 1 Data Cleaning

All missing values were imputed or dropped, as described above. Since our data has been validated by the Bureau of Transportation Statistics and Climate Data Online, outliers were investigated by reviewing the summary statistics of our data set.

```
[5]: dec_merged.describe()
```

```

[5]:      DAY_OF_MONTH  DAY_OF_WEEK  ORIGIN_AIRPORT_ID  DEP_DEL15  \
count  558026.000000  558026.000000      558026.000000  558026.000000
mean      15.830902      3.938745      12666.002996      0.208399
std       8.957760      2.085336      1514.187330      0.406164
min       1.000000      1.000000      10140.000000      0.000000
25%       8.000000      2.000000      11292.000000      0.000000
50%      16.000000      4.000000      12889.000000      0.000000
75%      23.000000      6.000000      13931.000000      0.000000
max      31.000000      7.000000      15919.000000      1.000000

      CANCELLED  CRS_ELAPSED_TIME  DISTANCE  DISTANCE_GROUP  \
count  558026.0      558026.000000  558026.000000  558026.000000
mean      0.0      148.552937      843.568687      3.844704
std      0.0      74.475448      604.827406      2.372199
min      0.0      34.000000      66.000000      1.000000
25%      0.0      94.000000      400.000000      2.000000
50%      0.0     130.000000      680.000000      3.000000
75%      0.0     179.000000     1075.000000      5.000000
max      0.0     705.000000     5095.000000     11.000000

      AIRLINE_ID  PILOTS_COPILOTS  PASSENGER_HANDLING  PASS_GEN_SVC_ADMIN  \
count  558026.000000      558026.000000      558026.000000      558026.000000
mean   19954.738880      6132.518447      7380.776432      9991.061352
std     368.971181      3163.783165      5905.764240      6417.203879
min    19393.000000      586.000000        0.000000      154.000000
25%    19790.000000     2444.000000     1407.000000     3592.000000
50%    19930.000000     7637.000000     8586.000000    15237.000000
75%    20314.000000     8989.000000     9668.000000    15502.000000
max    20436.000000     9293.000000    16888.000000    15809.000000

      MAINTENANCE      PRCP      SNOW      SNWD  \

```

count	558026.000000	558026.000000	558026.000000	558026.000000
mean	3576.673642	0.116641	0.048059	0.121935
std	3092.215270	0.352309	0.347030	0.806783
min	34.000000	0.000000	0.000000	0.000000
25%	898.000000	0.000000	0.000000	0.000000
50%	2482.000000	0.000000	0.000000	0.000000
75%	6122.000000	0.040000	0.000000	0.000000
max	9677.000000	7.130000	13.300000	18.100000

	TMAX	AWND
count	558026.000000	558026.000000
mean	56.160668	8.137934
std	14.612596	4.014022
min	8.000000	0.000000
25%	45.000000	4.920000
50%	56.000000	7.610000
75%	67.000000	10.510000
max	87.000000	25.720000

Evaluation of min and max values indicate that outliers do not exist. Also, thanks to the preprocessing of our datasets before merging, duplicates and formatting issues were already resolved.

## 1.1 1. Feature Selection

- We will disregard columns with IDs
- Also drop any redundant columns such as having only 1 distinct value
- Since we have both the distance and distance group column, it is redundant or DEP\_TIME\_BLK vs. ARR\_TIME\_BLK

```
[6]: dec_merged.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 558026 entries, 0 to 558025
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   DAY_OF_MONTH           558026 non-null int64
1   DAY_OF_WEEK            558026 non-null int64
2   OP_UNIQUE_CARRIER     558026 non-null object
3   TAIL_NUM                558026 non-null object
4   ORIGIN_AIRPORT_ID      558026 non-null int64
5   ORIGIN                  558026 non-null object
6   DEST                    558026 non-null object
7   DEP_DEL15              558026 non-null float64
8   DEP_TIME_BLK           558026 non-null object
9   ARR_TIME_BLK           558026 non-null object
10  CANCELLED               558026 non-null float64
11  CRS_ELAPSED_TIME        558026 non-null float64
```

12	DISTANCE	558026	non-null	float64
13	DISTANCE_GROUP	558026	non-null	int64
14	AIRLINE_ID	558026	non-null	int64
15	CARRIER_NAME	558026	non-null	object
16	PILOTS_COPILOTS	558026	non-null	int64
17	PASSENGER_HANDLING	558026	non-null	int64
18	PASS_GEN_SVC_ADMIN	558026	non-null	int64
19	MAINTENANCE	558026	non-null	int64
20	PRCP	558026	non-null	float64
21	SNOW	558026	non-null	float64
22	SNWD	558026	non-null	float64
23	TMAX	558026	non-null	float64
24	AWND	558026	non-null	float64

dtypes: float64(9), int64(9), object(7)

memory usage: 106.4+ MB

Of the 24 features remaining in our dataset, several were redundant upon inspection. `ORIGIN_AIRPORT_ID` is redundant with `ORIGIN`. `ARR_TIME_BLK`, which represents the arrival time block, is redundant with the combination of `DEP_TIME_BLK` and `CRS_ELAPSED_TIME`, or the schedule length of the flight, and `DISTANCE_GROUP`, or the distance of the flight. `AIRLINE_ID` is redundant with `CARRIER_NAME`. `DISTANCE` is redundant with the `DISTANCE_GROUP` binned feature. Since `CRS_ELAPSED_TIME` would be a function of the distance traveled, it was also dropped. `OP_UNIQUE_CARRIER` is redundant with `CARRIER_NAME`, which is also easier to interpret in subsequent sections, therefore `OP_UNIQUE_CARRIER` will be dropped. Since all cancelled flights were removed, the `CANCELLED` feature is irrelevant. This left 17 features in our dataset. Some features, such as `CARRIER_NAME`, `ORIGIN`, and `DEST` will not be used in modeling but are retained for later use.

```
[7]: dropped = ['TAIL_NUM', 'ORIGIN_AIRPORT_ID',
               ↪ 'ARR_TIME_BLK', 'CANCELLED', 'AIRLINE_ID', 'DISTANCE', 'CRS_ELAPSED_TIME',
               ↪ 'OP_UNIQUE_CARRIER']
dec_red = dec_merged.drop(dropped, axis=1)
dec_red.head()
```

```
[7]:   DAY_OF_MONTH  DAY_OF_WEEK  ORIGIN  DEST  DEP_DEL15  DEP_TIME_BLK  \
0             8             7    STL   SAN           0.0    1100-1159
1             8             7    STL   SAT           0.0    1200-1259
2             8             7    STL   SAT           0.0    2100-2159
3             8             7    STL   SEA           0.0    0900-0959
4             8             7    STL   SFO           1.0    1800-1859

   DISTANCE_GROUP  CARRIER_NAME  PILOTS_COPILOTS  \
0             7  Southwest Airlines Co.           8989
1             4  Southwest Airlines Co.           8989
2             4  Southwest Airlines Co.           8989
3             7  Southwest Airlines Co.           8989
4             7  Southwest Airlines Co.           8989
```

	PASSENGER_HANDLING	PASS_GEN_SVC_ADMIN	MAINTENANCE	PRCP	SNOW	SNWD	\
0	9668	15475	2482	0.02	0.0	0.0	
1	9668	15475	2482	0.02	0.0	0.0	
2	9668	15475	2482	0.02	0.0	0.0	
3	9668	15475	2482	0.02	0.0	0.0	
4	9668	15475	2482	0.02	0.0	0.0	

	TMAX	AWND
0	58.0	9.84
1	58.0	9.84
2	58.0	9.84
3	58.0	9.84
4	58.0	9.84

```
[8]: dec_red.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 558026 entries, 0 to 558025
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   DAY_OF_MONTH          558026 non-null int64
1   DAY_OF_WEEK           558026 non-null int64
2   ORIGIN                 558026 non-null object
3   DEST                  558026 non-null object
4   DEP_DEL15             558026 non-null float64
5   DEP_TIME_BLK          558026 non-null object
6   DISTANCE_GROUP        558026 non-null int64
7   CARRIER_NAME         558026 non-null object
8   PILOTS_COPILOTS       558026 non-null int64
9   PASSENGER_HANDLING    558026 non-null int64
10  PASS_GEN_SVC_ADMIN     558026 non-null int64
11  MAINTENANCE            558026 non-null int64
12  PRCP                   558026 non-null float64
13  SNOW                   558026 non-null float64
14  SNWD                   558026 non-null float64
15  TMAX                   558026 non-null float64
16  AWND                   558026 non-null float64
dtypes: float64(6), int64(7), object(4)
memory usage: 72.4+ MB
```

```
[9]: # Correlation Matrix for multicollinearity
plt.figure(figsize=(24, 12))
mask = np.triu(np.ones_like(dec_red.corr(), dtype=np.bool))
heatmap = sns.heatmap(dec_red.corr(), mask=mask, vmin=-1, vmax=1, annot=True)

# Make the full heat map visible
```

```

b, t = plt.ylim() # discover the values for bottom and top
b += 0.5 # Add 0.5 to the bottom
t -= 0.5 # Subtract 0.5 from the top
plt.ylim(b, t) # update the ylim(bottom, top) values
plt.show() # ta-da!


```

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:3:  
 DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To  
 silence this warning, use `bool` by itself. Doing this will not modify any  
 behavior and is safe. If you specifically wanted the numpy scalar type, use  
 `np.bool\_` here.

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

This is separate from the ipykernel package so we can avoid doing imports  
 until



## 1.2 2. Feature Creation

Since the day of the month did not show an easily identifiable pattern for departure delays, a new feature WEEK\_OF\_MONTH was created. This feature consists of 4 full (7-day) weeks and 1 partial (3-day) week. WEEK\_OF\_MONTH distributions for departure delay were evaluated to determine the optimal feature to use.

WEEK\_OF\_MONTH

1-7 = Week 1 = 1

8-14 = Week 2 = 2



15-21 = Week 3 = 3  
 22-28 = week 4 = 4  
 29-31 = partial week 5 = 5

While DEP\_TIME\_BLK showed- REDUCE CARDINALITY OF DEPARTURE TIME BLOCKS  
 (PREVIOUSLY 19 levels with inconsistent bucket size to 4 bins)

Redeye/Early Departure = 12:01 a.m. - 5:59 a.m. = 1  
 Morning Departure = 6:00 a.m - 11:59 a.m. = 2  
 Daytime Departure = 12:00 p.m. - 5:59 p.m = 3  
 Late Departure = 6:00 p.m. - 11:59 p.m. = 4

### 1.2.1 Create WEEK\_OF\_MONTH

```
[10]: def month_weeks_range(x):
        if x <= 7:
            return 1
        elif x <= 14:
            return 2
        elif x <= 21:
            return 3
        elif x <= 28:
            return 4
        else:
            return 5

dec_red['WEEK_OF_MONTH'] = dec_red['DAY_OF_MONTH'].apply(lambda x:
    ↪month_weeks_range(x))
dec_red['WEEK_OF_MONTH'].value_counts()
```

```
[10]: 3    129144
      1    128108
      2    125664
      4    122390
      5     52720
      Name: WEEK_OF_MONTH, dtype: int64
```

```
[11]: # Explore DAY OF MONTH with DEP_DEL15
Week = pd.crosstab(dec_red['WEEK_OF_MONTH'], dec_red['DEP_DEL15'])
Week['Total'] = Week.sum(axis=1)
Week.loc['Total'] = Week.sum()
Week['Percent_Delayed'] = ((Week.iloc[:,1])/((Week.iloc[:,0])+(Week.iloc[:,1])))
Week = Week.sort_values('Percent_Delayed')
Week
```

```
[11]: DEP_DEL15      0.0      1.0  Total  Percent_Delayed
WEEK_OF_MONTH
2          102803   22861  125664          0.181922
```

Total	441734	116292	558026	0.208399
3	101636	27508	129144	0.213003
1	100715	27393	128108	0.213827
4	95691	26699	122390	0.218147
5	40889	11831	52720	0.224412

```
[12]: # Drop DAY_OF_MONTH in place of WEEK_OF_MONTH
dec_red.drop(columns=['DAY_OF_MONTH'], inplace = True)
```

### 1.2.2 Transform DEP\_TIME\_BLK

```
[13]: dep_blk = {'0600-0659':2, '0700-0759':2, '0800-0859':2,
                '0900-0959':2, '1000-1059':2, '1100-1159':2,
                '1200-1259':3, '1300-1359':3, '1400-1459':3,
                '1500-1559':3, '1600-1659':3, '1700-1759':3,
                '1800-1859':4, '1900-1959':4, '2000-2059':4,
                '2100-2159':4, '2200-2259':4,
                '2200-2259':4, '2300-2359':4, '0001-0559':1}

dec_red['DEP_TIME_BLK'] = dec_red['DEP_TIME_BLK'].replace(dep_blk)
dec_red['DEP_TIME_BLK'].value_counts()
```

```
[13]: 2    213792
      3    197393
      4    131938
      1     14903
      Name: DEP_TIME_BLK, dtype: int64
```

```
[14]: # Explore DEP_TIME_BLK with DEP_DEL15
DEP = pd.crosstab(dec_red['DEP_TIME_BLK'], dec_red['DEP_DEL15'])
DEP['Total'] = DEP.sum(axis=1)
DEP.loc['Total'] = DEP.sum()
DEP['Percent_Delayed'] = ((DEP.iloc[:,1])/((DEP.iloc[:,0])+(DEP.iloc[:,1])))
DEP = DEP.sort_values('Percent_Delayed')
DEP
```

```
[14]: DEP_DEL15      0.0      1.0  Total  Percent_Delayed
DEP_TIME_BLK
1           13563      1340   14903           0.089915
2           183722     30070  213792           0.140651
Total       441734     116292  558026           0.208399
3           150278     47115  197393           0.238686
4            94171     37767  131938           0.286248
```

### 1.3 3. Feature Transformation

Since the non-numeric features remaining in our dataset will not be ingested in the model, but used later to enhance findings, interpretations and recommendations, and we are using the XGBoost Classification algorithm for modeling purposes, further transformations are unnecessary. XGBoost is not sensitive to transformations or scaling of features in the same way that decision trees and random forest are not. By not scaling our features, we remove the need to scale subsequent data ingested by the model and facilitate easier interpretability of our model with real-world data.

```
[15]: # List remaining features and types
dec_red.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 558026 entries, 0 to 558025
Data columns (total 17 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   DAY_OF_WEEK           558026 non-null int64
 1   ORIGIN                558026 non-null object
 2   DEST                 558026 non-null object
 3   DEP_DEL15            558026 non-null float64
 4   DEP_TIME_BLK         558026 non-null int64
 5   DISTANCE_GROUP       558026 non-null int64
 6   CARRIER_NAME        558026 non-null object
 7   PILOTS_COPILOTS     558026 non-null int64
 8   PASSENGER_HANDLING   558026 non-null int64
 9   PASS_GEN_SVC_ADMIN   558026 non-null int64
10   MAINTENANCE          558026 non-null int64
11   PRCP                 558026 non-null float64
12   SNOW                 558026 non-null float64
13   SNWD                 558026 non-null float64
14   TMAX                 558026 non-null float64
15   AWND                 558026 non-null float64
16   WEEK_OF_MONTH        558026 non-null int64
dtypes: float64(6), int64(8), object(3)
memory usage: 72.4+ MB
```

### 1.4 4. Balance the data set

```
[16]: # Identify the imbalanced class
dec_red['DEP_DEL15'].value_counts()
```

```
[16]: 0.0    441734
      1.0    116292
      Name: DEP_DEL15, dtype: int64
```

```
[17]: df_grouped_by = dec_red.groupby(["DEP_DEL15"])
```

```
df_balanced = df_grouped_by.apply(lambda x: x.sample(df_grouped_by.size().
    ↪min()))
df_balanced.reset_index(drop=True, inplace =True)
df_balanced
```

```
[17]:
```

	DAY_OF_WEEK	ORIGIN	DEST	DEP_DEL15	DEP_TIME_BLK	DISTANCE_GROUP	\
0	2	ATL	DEN	0.0	4	5	
1	6	MSP	FLL	0.0	2	6	
2	5	OGG	SFO	0.0	4	10	
3	5	IAD	SEA	0.0	3	10	
4	3	SEA	BOI	0.0	3	2	
...	...	...	...	...	...	...	
232579	2	SEA	BUR	1.0	2	4	
232580	1	PHX	TUS	1.0	4	1	
232581	4	SAN	SEA	1.0	3	5	
232582	6	TPA	ATL	1.0	3	2	
232583	6	DEN	IND	1.0	4	4	

	CARRIER_NAME	PILOTS_COPILOTS	PASSENGER_HANDLING	\
0	Southwest Airlines Co.	8989	9668	
1	Delta Air Lines Inc.	9293	15331	
2	United Air Lines Inc.	7637	16888	
3	Delta Air Lines Inc.	9293	15331	
4	Alaska Airlines Inc.	2893	1062	
...	...	...	...	
232579	Alaska Airlines Inc.	2893	1062	
232580	Mesa Airlines Inc.	1312	0	
232581	Delta Air Lines Inc.	9293	15331	
232582	Delta Air Lines Inc.	9293	15331	
232583	Southwest Airlines Co.	8989	9668	

	PASS_GEN_SVC_ADMIN	MAINTENANCE	PRCP	SNOW	SNWD	TMAX	AWND	\
0	15475	2482	0.00	0.0	0.0	56.0	11.86	
1	15809	6122	0.00	0.0	0.0	72.0	7.83	
2	15237	4991	0.39	0.0	0.0	81.0	7.38	
3	15809	6122	0.25	0.0	0.0	38.0	5.82	
4	5737	898	0.54	0.0	0.0	49.0	8.05	
...	...	...	...	...	...	...	...	
232579	5737	898	0.00	0.0	0.0	40.0	4.70	
232580	1205	34	0.00	0.0	0.0	62.0	6.49	
232581	15809	6122	0.00	0.0	0.0	69.0	2.91	
232582	15809	6122	0.32	0.0	0.0	80.0	10.07	
232583	15475	2482	0.18	2.8	1.2	29.0	19.24	

	WEEK_OF_MONTH
0	5
1	1

2	1
3	2
4	2
...	...
232579	4
232580	3
232581	1
232582	4
232583	4

[232584 rows x 17 columns]

```
[18]: # Confirm the majority class was undersampled
df_balanced['DEP_DEL15'].value_counts()
```

```
[18]: 0.0    116292
      1.0    116292
      Name: DEP_DEL15, dtype: int64
```

## 1.5 5. Split the data set

```
[19]: # partition model used features
model = df_balanced.drop(['ORIGIN', 'DEST', 'CARRIER_NAME'], axis = 1)

# Move DEP_DEL15 to first position for modeling
Dep = model['DEP_DEL15']
model.drop(labels=['DEP_DEL15'], axis=1, inplace = True)
model.insert(0, 'DEP_DEL15', Dep)
model
```

```
[19]:
```

	DEP_DEL15	DAY_OF_WEEK	DEP_TIME_BLK	DISTANCE_GROUP	PILOTS_COPILOTS	\
0	0.0	2	4	5	8989	
1	0.0	6	2	6	9293	
2	0.0	5	4	10	7637	
3	0.0	5	3	10	9293	
4	0.0	3	3	2	2893	
...	...	...	...	...	...	
232579	1.0	2	2	4	2893	
232580	1.0	1	4	1	1312	
232581	1.0	4	3	5	9293	
232582	1.0	6	3	2	9293	
232583	1.0	6	4	4	8989	

	PASSENGER_HANDLING	PASS_GEN_SVC_ADMIN	MAINTENANCE	PRCP	SNOW	SNWD	\
0	9668	15475	2482	0.00	0.0	0.0	
1	15331	15809	6122	0.00	0.0	0.0	
2	16888	15237	4991	0.39	0.0	0.0	

3	15331	15809	6122	0.25	0.0	0.0
4	1062	5737	898	0.54	0.0	0.0
...	...	...	...	...	...	...
232579	1062	5737	898	0.00	0.0	0.0
232580	0	1205	34	0.00	0.0	0.0
232581	15331	15809	6122	0.00	0.0	0.0
232582	15331	15809	6122	0.32	0.0	0.0
232583	9668	15475	2482	0.18	2.8	1.2

	TMAX	AWND	WEEK_OF_MONTH
0	56.0	11.86	5
1	72.0	7.83	1
2	81.0	7.38	1
3	38.0	5.82	2
4	49.0	8.05	2
...	...	...	...
232579	40.0	4.70	4
232580	62.0	6.49	3
232581	69.0	2.91	1
232582	80.0	10.07	4
232583	29.0	19.24	4

[232584 rows x 14 columns]

```
[20]: # Split all data into 90% train and 10% holdout
df_train, df_holdout = train_test_split(model, test_size=0.10,
↳stratify=model['DEP_DEL15'])

# Split holdout to 50% validation and 50% test
df_val, df_test = train_test_split(df_holdout, test_size=0.50,
↳stratify=df_holdout['DEP_DEL15'])
```

add code to upload to bucket

```
[21]: df_train.shape
```

```
[21]: (209325, 14)
```

```
[22]: df_val.shape
```

```
[22]: (11629, 14)
```

```
[23]: df_test.shape
```

```
[23]: (11630, 14)
```

```
[30]: # Save train data set for modeling to model_data folder in bucket
csv_buffer=StringIO()
df_train.to_csv(csv_buffer, index=False)
```

```
BUCKET_NAME = 'ads-508-airline'
FileName= 'model_data/df_train.csv'
```

```
s3csv = boto3.client('s3')
```

```
response=s3csv.put_object(Body=csv_buffer.getvalue(),
                           Bucket=BUCKET_NAME,
                           Key=FileName)
```

```
[31]: # Save validation data set for modeling to model_data folder in bucket
csv_buffer=StringIO()
df_val.to_csv(csv_buffer, index=False)
```

```
BUCKET_NAME = 'ads-508-airline'
FileName= 'model_data/df_val.csv'
```

```
s3csv = boto3.client('s3')
```

```
response=s3csv.put_object(Body=csv_buffer.getvalue(),
                           Bucket=BUCKET_NAME,
                           Key=FileName)
```

```
[32]: # Save test data set for modeling to model_data folder in bucket
csv_buffer=StringIO()
df_test.to_csv(csv_buffer, index=False)
```

```
BUCKET_NAME = 'ads-508-airline'
FileName= 'model_data/df_test.csv'
```

```
s3csv = boto3.client('s3')
```

```
response=s3csv.put_object(Body=csv_buffer.getvalue(),
                           Bucket=BUCKET_NAME,
                           Key=FileName)
```

```
[ ]:
```